Build a personal Jarvis: memory, routing, voice and explicit tool access

Architecture before product shopping

Both sources describe an assistant with a persistent home, an execution harness, a model, memory, skills, channels and an interface. They are alternative implementations, not a single universal install script. The Chase AI workflow emphasizes Claude Code/Codex, Obsidian, local speech and typed routing; the nine-step workflow emphasizes Hermes Agent, Honcho and optional hosted voice.

Choose a route

The creator's Agentic OS post (https://www.chaseai.io/blog/build-claude-code-agentic-os-3-steps) places the complete package in a paid community. The source research could not establish a free official V2 repository. Check today's offer directly; old promotional prices and expiry dates in the individual guide are historical.

For a modular route, start with the official Hermes quickstart (https://hermes-agent.nousresearch.com/docs/getting-started/quickstart/) and Honcho repository (https://github.com/plastic-labs/honcho). Select supported versions and the desired hosted/self-hosted memory mode before using their installation commands. An "open-source harness" does not make hosted models or managed memory free.

Build in a testable order

1. Run a text-only fixture task and log its tool calls and outcome.
2. Add durable memory with explicit read/write scope and recovery. Obsidian's Markdown vault and Honcho's memory service have different storage and integration contracts; choose one first.
3. Add a model/router adapter only if a measured workload benefits. Jev returns typed decisions, not a complete conversation; preserve a deterministic fallback when it is unavailable.
4. Add local speech using whisper.cpp (https://github.com/ggml-org/whisper.cpp) and Kokoro (https://github.com/hexgrad/kokoro), or a hosted service after checking current commercial terms and latency. Local speech does not make the downstream LLM local.
5. Add one read-only calendar or document connector. Expand scopes only when the task needs them and the integration has been tested.
6. Add the dashboard last. A polished UI is not evidence of reliable actions.

Important corrections and limits

Codex is an agent product with model choices, not a standalone interchangeable foundation model. Credentials belong outside the vault. Background availability needs a machine, process management and a recovery plan. An iMessage bridge is not a generic Apple messaging API; it introduces its own machine and account requirements. The preserved source guides contain dated pricing/model examples, which must be rechecked before purchasing.

Project brief

Build a personal research assistant over fixture Markdown files, with one task inbox and one scoped read-only calendar adapter. Require task evidence, a bounded tool allowlist and a visible confirmation path for future write actions. Test permission denial, stale memory, model timeout and duplicate requests before adding voice.

```
Implement a text-first personal assistant with explicit memory and tool
contracts. Use fixtures and a mock model initially. Add one documented model
adapter, then optional speech. Log outcomes, expose unresolved work, and test
recovery and scope denial. Do not connect personal accounts automatically.
```

The shared-memory vault and Jev guides provide focused expansions. Both original implementations and all original source links remain in this topic's Files and Sources.

Original sources

- https://www.instagram.com/p/DdpdeD-HPsJ/?img_index=6
- https://www.instagram.com/reel/DdrQGLcT-3O/