AI model-name comedy: distinguish real models, derivatives and invented labels

Researched 2026-10-20.

Read this as comedy

The Muse dossier identifies a deadpan skit that mixes real model terminology with invented or garbled names. It contains no usable installation demonstration. An auto-transcript can also turn a joke into a plausible URL, so the spoken names are not a shopping list.

The real concepts behind the names

Microsoft's Phi-4 model card (https://huggingface.co/microsoft/phi-4) documents a real base model. Heretic's upstream repository (https://github.com/p-e-w/heretic) describes automated removal of refusal behavior through model modification. These sources establish those concepts; they do not authenticate every model label, quantization, performance claim or invented website in the reel.

A reduction in refusals is different from stronger reasoning, factual accuracy or better coding. A modified checkpoint can perform differently from its base model and has its own provenance and license obligations. This guide does not recommend a particular derivative or repeat the dossier's unverified free-gateway model IDs and rate limits.

A practical verification checklist

Before evaluating any named model, locate a maintainer-owned repository or publisher model card. Record the exact repository, revision, base model, license, intended use, file format and expected memory requirement. Separate the model author from a quantization uploader. Check whether a benchmark describes the same revision and inference configuration you would run.

If the only evidence is a comedy transcript, label the name unresolved. Do not invent a download link by combining fragments of the audio. No need to visit a guessed domain just to prove it does not exist.

What to compare

Use ordinary benign tasks with known answers: a short coding exercise with tests, a document extraction fixture and a factual question set where unsupported answers are counted. Keep the same prompt, context and sampling configuration for both checkpoints. Record correctness, latency, memory use and abstention quality. A refusal count alone does not measure useful capability.

This is an evaluation plan, not a claim that a small local model beats a frontier model or runs on a specified laptop. Actual hardware requirements depend on parameter count, precision, context and runtime.

Reusable prompt

```
Turn this model-name list into an evidence table.
For each entry, distinguish an official release, a third-party derivative,
a transcription uncertainty and a fictional name.
Use publisher sources and model cards; do not guess domains or downloads.
Propose benign repeatable tasks and compare accuracy, latency and resources.
Keep unavailable measurements blank and make no installation changes.
```

Optional project: model provenance cards

Build a small static catalog whose records contain source URL, revision, base model, license, format, verification date and an evidence note. Add a field for "unresolved" so the interface does not force a fake product match. The first version only needs local JSON, filtering and links to original model cards.

Limits and next action

The original transcript remains attributed to Muse and may contain mishears. This research did not run a model, modify weights, install a runtime or validate the skit's invented names. Next: use the provenance checklist on one real publisher model rather than chasing the entire joke list.

Original sources

- https://www.instagram.com/reel/DeevahVsCTK/